

Designing Constitutional Review Institutions by Simulation

Anonymous Author(s)

Abstract

Constitutional review institutions face a recurring design problem: they need enough independence to preserve rights and legal continuity, but enough accountability to avoid becoming unreviewable partisan veto points. We introduce a computational design-search model for comparing constitutional-review institutions under common docket and political assumptions. The model represents appointment systems, court size, tenure, removal, recusal, emergency procedure, voting thresholds, panel and en banc review, cross-checking courts, constitutional councils, legislative overrides, admissibility filters, lower-court splits, claimant type, bar capital, claim strength, vehicle quality, repeat-player advantages, forum shopping, enforcement capacity, and post-decision political response. The current campaign evaluates twenty-six institutional designs across twenty ordinary and adversarial assumption cases. A filed-universe intake model routes cases through review paths, court-requested-response and CVSG-style screening, and admission filters before merits review, and the reports keep direct outputs such as rights-claimant success, public confidence, lower-court compliance, government noncompliance, precedent durability, emergency downstream effects, and doctrinal depth separate from derived indices. The results do not identify an optimal court. They show how design bundles shift tradeoffs among legal stability, rights protection, emergency-docket abuse, partisan alignment, legitimacy, constitutional conflict, and democratic responsiveness.

Keywords: constitutional review, judicial politics, institutional design, simulation, courts

1 Design Question

The central question is not whether one court design is ideal in isolation. It is how bundles of constitutional-review rules alter institutional tradeoffs when they face the same filed universe, the same political environment, and the same strategic pressures. A court can be independent while losing democratic legitimacy, responsive while failing rights protection, stable while entrenching poor doctrine, or active while creating constitutional conflict. The simulator is built to expose those tensions rather than collapse them into a single institutional label.

This makes the paper a computational institutional-design exercise for judicial politics and comparative public law. The simulator is not intended to forecast any real supreme court. It is a design-search scaffold for asking how appointment regularity, emergency procedure, admission filtering, voting thresholds, override rules, council warnings, cross-court checks, and post-decision compliance strategies change the distribution of outcomes. That framing is especially important because many proposed court reforms are defended one rule at a time, even though their effects depend on the surrounding institutional bundle.

2 Contribution

The paper makes three contributions. First, it treats constitutional-review reform as a design-bundle problem rather than as a sequence of isolated proposals. Court size, tenure, appointment rules,

emergency procedure, voting thresholds, review access, and override rules interact, so we compare bundles under shared assumptions instead of ranking individual reforms by intuition. Second, it separates direct model outputs from derived indices. The campaign reports legal stability, rights protection, rights-claimant success, domain-specific claimant success, emergency grant behavior, public confidence, elite acceptance, lower-court compliance, and constitutional conflict before any summary score is used. Third, it provides a reproducible design-search workflow that links scenario definitions, normalized calibration inputs, deterministic seeds, generated reports, and manuscript figures.

These contributions are methodological rather than doctrinal. We do not argue that simulation can settle constitutional design. We argue that explicit simulation can make institutional tradeoffs more inspectable, especially when reforms that look attractive in isolation shift costs into emergency procedure, admission filtering, public legitimacy, interbranch conflict, or post-decision compliance.

3 Claim Boundaries and Interpretation

The manuscript separates empirical claims, synthetic findings, and speculative design recommendations. Empirical claims are limited to source descriptions, normalized calibration guardrails, cited literature, and reproducible project artifacts. They support statements such as which source families are used for plausibility checks, what institutional features are represented in the code, and whether generated artifacts match their manifests. They do not establish causal effects of any reform in the world.

Synthetic findings are outputs of the simulator under stated coding choices, deterministic seeds, and assumption cases. Statements about a scenario reducing shadow-docket abuse, increasing administrative cost, changing lower-court compliance, or moving public confidence are therefore simulation findings, not empirical findings about observed courts. The generated tables and figures are best read as mechanism diagnostics: they show where costs move when a rule bundle changes.

Speculative design recommendations are conditional inferences from those synthetic findings. The paper does not recommend weakening or strengthening courts as an end in itself. The design objective used here is narrower: reduce arbitrary emergency power while preserving rights protection, visible legal reasoning, lower-court settlement, and credible implementation. Reforms that reduce emergency opacity while increasing noncompliance, recusal pressure, or rights failure are not treated as improvements merely because they restrain a court.

4 Theory and Design Space

The model joins three law-and-courts traditions that are often studied separately. The first treats judges as purposive actors whose choices depend on legal goals, policy preferences, colleagues, litigants, and other institutions (Segal and Spaeth, 2002; Epstein and Knight, 1998; Clayton and Gillman, 1999). The second treats constitutional review as an institutional settlement among political actors who face uncertainty, seek insurance against future exclusion, or try to preserve influence through courts (Ginsburg, 2003; Hirschl, 2004; Dixon and Ginsburg, 2017). The third asks how courts preserve authority when they depend on elected officials, lower courts, publics, and repeat players for implementation (Vanberg, 2005; Staton, 2010). A useful simulator for constitutional-review design has to represent all three: judicial choice, institutional allocation, and post-decision compliance.

Comparative constitutional design also makes clear that “judicial review” is not one institution. Strong-form review, weak-form review, pre-enactment review, legislative override, individual

complaint, amparo, concrete referral, constitutional council screening, specialized constitutional courts, and ordinary supreme courts expose different actors to different veto points (Tushnet, 2008; Garbaud, 2013; Sweet, 2000; Brewer-Carías, 2009; Kommers and Miller, 2012). Override and weak-form review scholarship also shows that post-decision political response is part of the institution rather than an afterthought (Hogg and Bushell, 1997; Comella, 2009). That is why this project models access path, review timing, remedy threshold, emergency procedure, and override rule as separate design fields. The relevant unit of comparison is the bundle, not a single reform label.

Emergency procedure is treated separately because it changes both timing and legitimacy. The literature on the U.S. shadow docket emphasizes that orders outside the ordinary merits process can be consequential precisely because they are faster, less reasoned, and less procedurally regular (Baude, 2015; Vladeck, 2023; Kastellec and Taboni, 2026). The simulator therefore separates ordinary merits review from emergency applications, shadow relief, public disagreement, reasons, merits follow-through, and status-quo effects. This makes emergency legality, rights protection, and legitimacy costs visible instead of burying them in an aggregate activity rate.

The design problem also differs from ordinary policy optimization. Constitutional-review institutions produce several goods that can conflict: rights protection, legal stability, responsiveness to democratic majorities, settlement of interbranch conflict, and visible decisional legitimacy. A design that improves one good can undermine another by changing access, delay, information, coalition structure, or implementation incentives. For that reason, the simulator does not search for a single welfare-maximizing court. It asks whether a design bundle moves costs into predictable locations and whether those costs remain visible in direct outputs.

This framing matters for reform debates. A proposal to impose term limits is not only a tenure proposal; it changes vacancy timing, appointment incentives, and the value of partisan capture. A proposal to require reasoned emergency orders is not only a transparency proposal; it changes speed, emergency grant behavior, public disagreement, and merits follow-through. A proposal to add a constitutional council or judicial council is not only a rights-screening proposal; it changes *ex ante* review, appointment insulation, later judicial workload, and the political salience of invalidation (Garoupa and Ginsburg, 2009). The simulation architecture therefore treats legal form, access path, and political response as jointly produced institutional behavior.

5 Expectations

The simulations are exploratory, but they are not atheoretical. The campaign is organized around five design expectations that can be evaluated in direct outputs rather than only in a final score.

First, regularized appointments should reduce incentives to manipulate appointment timing and should soften public-legitimacy penalties, but the effect should be modest unless the regularized system also changes confirmation control or partisan appointment opportunities. Term limits are therefore expected to improve public confidence and reduce strategic pressure more than they transform rights outcomes.

Second, fragmented or screened appointment systems should reduce partisan alignment and some appointment-manipulation pressure, but they can add administrative cost and may not improve rights protection by themselves. The model therefore expects appointment screening to move legitimacy and alignment more than merits outcomes.

Third, emergency relief without merits follow-through should be a distinct legitimacy risk. Designs that require merits follow-through, reason giving, or full-court referral should reduce shadow-docket abuse and emergency legitimacy risk, but may pay for that improvement through more accelerated merits work or slower emergency handling.

Fourth, panel routing, cross-checking courts, dual supreme courts, and constitutional councils should not be treated as simple activity reducers. They may improve error correction, ex ante screening, or cross-institutional moderation, but they can also raise administrative cost, disagreement, or constitutional conflict if the rules for resolving institutional disagreement are weak.

Fifth, legislative override rules and accountability devices should increase democratic responsiveness only when they do not collapse into repeated override loops, rights carveout evasion, or symbolic compliance. The relevant output is therefore not only whether override is available, but whether it changes practical implementation, rights protection, elite acceptance, and constitutional conflict.

6 Model

Each generated world contains a pool of potential justices, a filed universe of constitutional disputes, and a legislative-output profile. Filed cases vary by legal ambiguity, rights burden, claim strength, vehicle quality, claimant type, bar capital, democratic mandate, partisan salience, emergency pressure, requested emergency relief, legislative quality, constitutional-conflict potential, public attention, override pressure, lower-court conflict, lower-court error risk, lower-court split depth, genuine lower-court split status, lower-court ideological drift, lower-court resistance risk, certiorari pressure, forum-shopping pressure, pre-review settlement pressure, strategic plaintiff selection, repeat-player advantage, government noncompliance risk, enforcement capacity, emergency opportunism, and recusal incentive pressure. We separate facial challenges, as-applied challenges, election disputes, emergency stay applications, executive-power disputes, administrative-law challenges, structural cases, economic-regulation cases, and rights claims. Legal-domain profiles shift baseline case attributes so rights, election, executive-power, administrative, structural, and economic-regulation cases do not behave as interchangeable labels.

A review-path layer distinguishes paid certiorari, IFP certiorari, discretionary certiorari, direct constitutional complaint, amparo, filtered QPC-style referral, abstract ex ante and ex post review, concrete referral, emergency application, interbranch dispute, and electoral review. Admission filtering turns filed disputes into screened-out matters or merits transfers using petition type, court-requested response, CVSG request, SG signal, amicus intensity, lower-court split depth, genuine split status, split maturity, relists, claimant type, bar capital, claim strength, vehicle quality, specialist counsel, vehicle defects, forum-shopping pressure, settlement pressure, repeat-player advantages, strategic plaintiff selection, and conditional reversal probability. There is no preliminary fixed-size top-case selector before admission; the modeled merits docket is an output of the intake process. The state transition is filed matter $\rightarrow$ petition or referral $\rightarrow$ response or CVSG screening when applicable $\rightarrow$ admission decision $\rightarrow$ settlement, merits transfer, emergency order, or screen-out $\rightarrow$ post-decision response.

Each scenario supplies a court design. The design selects reviewers, applies recusal rules, routes cases to full-court or panel review, determines emergency relief procedure, applies voting thresholds, optionally invokes cross-checking or council review, and optionally allows legislative or popular override. Emergency applications carry applicant type, application class, requested relief, response request, full-court referral, status-quo effect, public disagreement risk, and merits follow-through. Emergency orders can also create downstream effects through status-quo disruption, opportunistic emergency filing, unexplained relief, merits follow-through, and later public disagreement. Strategic actors choose whether to comply, evade, reenact, flood emergency applications, organize override campaigns, or pressure appointments. Their response is split into formal legal response and practical implementation response, including repeal, replacement, narrowed reenactment, weak-form override,

amendment, court-curbing, delay, administrative substitution, symbolic compliance, bureaucratic resistance, and open noncompliance.

The model uses shared worlds to make institutional comparisons sharper. Within a run, every scenario sees the same filed matters and legislative-output profile. Differences in outputs therefore come from the institutional rule bundle and its induced strategic responses, not from a different case draw. Across runs and stress cases, the world generator changes appointment capture, emergency pressure, rights risk, democratic mandate, public attention, executive defiance, and imported legislative quality. This gives the campaign both within-world comparison and across-assumption sensitivity.

The current design surface is deliberately broader than the tabled scenarios. Court size may be odd or even; terms may be life, fixed, staggered, renewable, or retirement-age bounded; removal may be impeachment-like, supermajoritarian, or external-review based; recusal may leave replacement, vacancy, or quorum consequences; and remedies may require different thresholds for invalidation, emergency relief, or override. The scenario catalog includes randomized merits panels with en banc correction, mandatory written emergency reasoning, automatic merits follow-up, strong recusal enforcement, jurisdiction-stripping constraints, legislative override windows, constitutional remand, public-interest litigation filtering, and bundled systems that combine emergency integrity, remand-and-override, panel-and-jurisdiction, or council-and-concrete-review safeguards. The diagnostics vary assumptions around these bundles rather than presenting every possible combination as if it were substantively meaningful.

That choice keeps the paper focused on institutional mechanisms rather than arbitrary parameter permutations.

An individual simulated dispute moves through the model as follows. First, the filed universe draws a legal domain and case type, then applies the imported legislative-output profile to the case's quality, mandate, rights burden, and conflict potential. A weakly supported, high-salience statute with a high rights burden is therefore not simply "a rights case"; it is a rights case with a particular legislative pedigree, claimant profile, legal vehicle, and public environment. Second, the access-path module assigns the procedural route. A similar substantive dispute may enter as certiorari, constitutional complaint, amparo, concrete referral, abstract review, emergency application, or interbranch dispute. Third, the admission filter applies path-specific signals. Lower-court conflict, lower-court split depth, genuine split status, lower-court ideological drift, forum shopping, SG support, court-requested response, CVSG request, amicus intensity, split maturity, relists, bar capital, claim strength, vehicle quality, repeat-player advantage, strategic plaintiff selection, counsel quality, vehicle defects, settlement pressure, and conditional reversal probability affect cert-like paths, while complaint and referral paths face different gatekeeping. Fourth, admitted disputes can settle before review, move to merits review, or receive emergency disposition under the scenario's court design. The design determines reviewers, recusals, quorum effects, panel or en banc routing, voting thresholds, reason giving, cross-court or council checks, and emergency merits follow-through. Fifth, the decision creates a formal legal response, a practical implementation response, a precedent-durability value, and an emergency downstream effect. A legislature or executive may repeal, replace, narrow, override, delay, symbolically comply, administratively substitute, or openly defy; enforcement capacity, lower-court resistance, lower-court compliance, and elite acceptance then feed public confidence and constitutional conflict.

The variables in that walkthrough are not fitted coefficients. They are explicit coding choices disciplined by source-range smoke tests, shared-world comparisons, and sensitivity sweeps. This is why the paper reports direct outputs before any score: a design can improve public confidence by reducing opaque emergency relief while leaving domain-specific claimant success nearly unchanged, or it can raise democratic responsiveness while also increasing override loops and conflict. The

method is useful only if those movements remain visible rather than being hidden inside a single index.

7 Calibration and Source-Range Checks

The calibration layer is deliberately framed as a source-range smoke test rather than as a claim of full validation. Normalized source observations are drawn from the Supreme Court Database, Harvard Law Review Supreme Court statistics, the Supreme Court Shadow Docket Database, and Black and Epstein’s recusal data (Supreme Court Database, 2025; Harvard Law Review, 2025; Kastellec and Taboni, 2026; Black and Epstein, 2005). Comparative synthesis rows generated during project research are retained as structured design-context inputs with source names, URLs, confidence labels, validation-use fields, denominator notes, and coverage scopes, but the synthesis memos themselves are not treated as empirical authorities. Newer research tables on certiorari screening, emergency applications, comparative access paths, lower-court alignment, implementation delay, claimant type, and repeat-player effects are imported into a separate source-register directory with strict-validation, loose-calibration, proxy-context, and design-prior labels. A row-level provenance manifest records the normalized file, source version, source URL, raw-source requirement, transformation script, extraction mode, validation use, and denominator specification for each calibration family. The current calibration run reports classified source-range smoke tests rather than validation claims; the targets are explicitly classified as strict validation checks, loose calibration checks, proxy sanity checks, mechanism sanity checks, or model-prior checks.

The model distinguishes three categories that should not be blurred. Observed inputs anchor source ranges or qualitative rules. Coding decisions translate legal institutions into parameterized mechanisms. Simulated outputs are exploratory comparisons under those mechanisms. Several quantities remain intentionally loose: SCOTUSblog significant-application counts are not full emergency denominators; same-term grant and docket-flow ratios are not closed petition cohorts; settlement-before-review lacks a clean cross-jurisdiction validation series; raw cross-national invalidation rates are not directly comparable without institutional context; and source tables describing appointment actors or court structure are not treated as direct measures of independence. The generated source-range appendix reports raw source minima, percentiles, medians, maxima, terms, and source keys, while the campaign reports preserve deterministic seeds and scenario definitions for reproduction.

Table 1 reports representative guardrails rather than every source row. It is meant to show the type of check being applied: a current-like design should have plausible docket composition, invalidation, recusal, and emergency-order behavior, while a no-emergency-relief-without-merits design should sharply suppress shadow relief without erasing emergency merits acceleration. “Within range” does not mean the model is causally identified or externally validated. It means the run has not drifted outside ranges that would make the comparison facially implausible for the stated guardrail use.

Table 2 is a safeguard against overstating the evidence. Strict-validation rows are closest to empirical denominators, loose-calibration rows are plausibility checks, proxy rows are warning lights, mechanism rows test whether coded mechanisms move in the expected direction, and model-prior rows document institutional coding assumptions. A scenario does not become empirically validated because it passes a proxy or mechanism row.

Table 3 is the denominator audit. Certiorari screening, emergency applications, individual complaints, QPC/concrete referrals, lower-court compliance, and override/remand behavior are shown as separate source-comparison pathways. The table records both the manuscript-use label

Table 1: Selected calibration guardrails

Guardrail	Scenario	Use	Evidence	Observed	Range	Source basis	Status
Merits transfer	Current-like	model prior check	model prior	0.387	0.250–0.850	source register; model fallback	within
Paid CFR stage	Current-like	loose calibration	research synthesis	0.056	0.027–0.067	research tables; source range	within
IFP CFR stage	Current-like	loose calibration	research synthesis	0.032	0.000–0.040	research tables; source range	within
Invalidation	Current-like	strict validation	raw or primary summary	0.093	0.000–0.223	SCDB; source range	within
Emergency filings	Current-like	proxy sanity check	raw or primary summary	0.046	0.000–0.223	shadow database; source range	within
Shadow abuse	Current-like	loose calibration	raw or primary summary	0.204	0.000–0.657	shadow database; source range	within
Partisan alignment	Commission	proxy sanity check	raw or primary summary	0.020	0.000–0.657	shadow database; source range	within
Emergency restraint	No-gap	mechanism sanity check	raw or primary summary	0.007	0.000–0.657	shadow database; source range	within

Table 2: Calibration-use classification summary

Guardrail use	Rows	Within range	Interpretation
loose calibration	10	10	plausibility range, not point validation
mechanism sanity check	2	2	context row
model prior check	2	2	context row
proxy sanity check	4	4	context row
strict validation	4	4	closest to source-denominator checks

and the denominator compatibility judgment, because a source row that is strict within its own dataset can still be only proxy context for a simulator metric. This is intentionally stricter than the calibration table: a metric may be useful as a synthetic diagnostic while still needing a direct primary-source validation target.

8 Metrics

The headline metrics are legal stability, rights protection, partisan alignment, shadow-docket abuse, legitimacy, reversal rate, constitutional conflict, and democratic responsiveness. Legal stability is decomposed into precedent stability, statutory stability, interbranch compliance, precedent durability, and lower-court compliance. Emergency legitimacy risk is reported separately from shadow-docket abuse and emergency downstream effect, so opacity, procedure legitimacy, and later disruption are not collapsed into the same quantity as abusive merits-like relief. Rights-claimant success is reported both as an aggregate direct output and by legal domain for rights, structural, election, executive-power, administrative-law, and economic-regulation cases. Intake diagnostics now also report court-requested response, CVSG request, paid and IFP CFR rates, claimant-type composition, bar capital, claim strength, vehicle quality, and genuine lower-court split rate. Actor-learning diagnostics report repeat-player learning, emergency-incentive learning, and compliance learning as derived measures, while the process state itself updates those incentives across repeated decisions.

The reported directional score is only a reading aid. It averages a stability/rights score, a

Table 3: Pathway-specific denominator audit

Pathway	Construct	Sim. metric	Current	Source metric	Use	Compatibility	Next action
certiorari	paid certiorari intake share	paid cert petition share	0.235	paid petition share	loose calibration	near pathway match	expand term coverage
certiorari	paid court-requested response	paid cfr request rate	0.061	cfr rate paid	loose calibration	near pathway match	expand term coverage
certiorari	grant/admission funnel	certiorari admission rate	0.195	us scotus plenary review rate all docketed	proxy context	denominator mismatch	do not treat as validation; seek direct denominator before causal claims
emergency	emergency application presence	emergency stay docket rate	0.092	emergency stay docket rate	proxy context	denominator mismatch	do not treat as validation; seek direct denominator before causal claims
emergency	emergency relief grants	emergency grant conditional rate	0.242	noncapital grant rate overall	loose calibration	conditional near match	expand term coverage
complaint referral	individual complaint admission	admission rate	0.533	spain amparo admission rate	proxy context	denominator mismatch	do not treat as validation; seek direct denominator before causal claims
lower court compliance	lower-court alignment	lower court compliance	0.514	district court alignment shock same direction percentage points	proxy context	mechanism proxy	do not treat as validation; seek direct denominator before causal claims
lower court compliance	government noncompliance	government noncompliance rate	0.181	federal agency narrow compliance share	loose calibration	mechanism proxy	do not treat as validation; seek direct denominator before causal claims
complaint referral	constitutional remand/deferred remedy	constitutional remand rate	0.000	fr qpc delayed effect rate decided qpc	proxy context	mechanism proxy	do not treat as validation; seek direct denominator before causal claims
override remand	override pressure after invalidation	override attempt rate	0.000	invalidation rate	proxy context	mechanism proxy	do not treat as validation; seek direct denominator before causal claims

legitimacy/control score, rights-claimant success, precedent durability, lower-court compliance, elite acceptance, and administrative feasibility. Direct outputs such as certiorari admission, court-requested response, CVSG request, bar capital, claim strength, vehicle quality, genuine lower-court split rate, lower-court split depth, strategic plaintiff selection, repeat-player advantage, government noncompliance, recusal incentive pressure, emergency downstream effect, rights-claimant success, doctrinal depth, remedial breadth, lower-court compliance, elite acceptance, public confidence, and coalition fragmentation are reported separately and should not be treated as interchangeable components of one welfare function. Appendix Table 11 records the denominator or scale for representative metrics so empirical source checks, synthetic outputs, and score-reading aids are not conflated.

Figure 1 illustrates why the domain-specific claimant-success measures are kept separate. Designs can look similar on aggregate rights protection while moving success rates differently for rights, structural, election, executive-power, administrative-law, and economic-regulation disputes.

9 Campaign Results

The current campaign compares twenty-six institutional scenarios across twenty assumption cases: the original broad cases, sensitivity sweeps for appointment capture, emergency pressure, rights risk, and democratic mandate, plus adversarial stress cases for appointment timing, emergency-application flooding, override evasion, recusal pressure, and court-expansion retaliation. The reported run used 80 runs per case, 64 filed matters per run, seed 20260501, and an imported legislative-output profile from a separate legislative simulator report contract.

Figure 2 plots public confidence against constitutional conflict for selected designs, with marker shade representing shadow-docket abuse. The figure is diagnostic rather than normative: movement toward lower conflict and higher confidence is usually attractive, but the figure intentionally leaves

	Rights	Struct.	Elect.	Exec.	Admin.	Econ.
Current-like	0.41	0.42	0.96	0.34	0.41	0.22
18-year terms	0.38	0.41	0.74	0.22	0.32	0.20
Written reasons	0.36	0.41	0.73	0.23	0.29	0.19
Auto merits	0.47	0.49	0.96	0.34	0.47	0.25
No merits gap	0.46	0.45	0.94	0.32	0.44	0.23
Recusal enforce.	0.39	0.40	0.73	0.23	0.34	0.20
Random panels	0.38	0.40	0.74	0.21	0.32	0.19
Integrity pkg.	0.47	0.46	0.95	0.33	0.42	0.24

Domain-specific claimant-success rates; darker cells are higher.

Figure 1: Domain-specific rights-claimant success for selected institutional designs. The aggregate rights-protection score can hide domain shifts, especially for election, structural, executive-power, administrative-law, and economic-regulation cases.

rights protection and democratic responsiveness outside the two-dimensional display.




Figure 2: Public-confidence and constitutional-conflict tradeoff for selected institutional designs. Marker shade increases with shadow-docket abuse, so the figure separates legitimacy and conflict from emergency-docket opacity instead of folding them into one score.

Table 4 should be read as a clustered comparison, not as a tournament table. Differences below .010 in the directional score are treated as substantively indistinguishable, and even larger differences should be interpreted through the direct outputs before drawing any design conclusion.

Table 4: Selected v2 campaign averages, grouped by emergency-power and rights-protection profile

Scenario	Rights	Shadow	Emerg. risk	Downstream	Gov. noncomp.	Lower ct.	Prec. dur.	Public	Score
Stylized current U.S.-like court	0.634	0.306	0.314	0.186	0.181	0.514	0.599	0.344	0.552
18-year staggered terms	0.620	0.105	0.235	0.112	0.147	0.536	0.633	0.488	0.565
Mandatory written emergency reasoning	0.618	0.076	0.245	0.128	0.149	0.535	0.649	0.489	0.562
Automatic merits follow-up	0.631	0.030	0.221	0.089	0.148	0.542	0.588	0.545	0.560
No emergency relief without merits review	0.632	0.022	0.220	0.088	0.146	0.544	0.605	0.571	0.566
Independent recusal	0.621	0.106	0.238	0.113	0.148	0.536	0.632	0.519	0.559
enforcement with substitutes									
Randomized merits panels	0.622	0.107	0.237	0.114	0.146	0.536	0.639	0.543	0.560
with en banc correction									
Emergency integrity package	0.633	0.030	0.222	0.089	0.146	0.543	0.598	0.575	0.559
Constitutional remand before invalidation	0.619	0.106	0.238	0.113	0.145	0.554	0.698	0.541	0.566
Remand with override window	0.616	0.077	0.247	0.128	0.146	0.553	0.711	0.540	0.563
Public-interest litigation filter	0.626	0.106	0.238	0.112	0.147	0.537	0.643	0.509	0.563

The stable front-line cluster contains several emergency-procedure and process-regularity reforms that preserve rights protection while lowering shadow-docket abuse or emergency downstream effects. Mandatory written emergency reasoning, automatic merits follow-up, no-relief-without-merits review, strong recusal enforcement, randomized panels with en banc correction, constitutional remand, and public-interest filters are therefore compared as different ways to reduce arbitrary emergency power rather than as a simple ladder from weak courts to strong courts. The stylized current U.S.-like court can have substantial rights protection in this run while still producing higher shadow-docket abuse and emergency legitimacy risk. That is the central tradeoff, not a claim that one directional score wins.

Table 5: Litigation-pipeline and downstream-enforcement diagnostics

Scenario	Cert admit	CFR	CVSG	Bar	Claim	Vehicle	Genuine split	Split	Resistance	Forum	Settlement	Plaintiff sel.	Repeat player	Enforcement	Emerg. opp.	Recusal press.	Gov. noncomp.	Emerg. down.	Prec. dur.
Current-like	0.195	0.118	0.001	0.570	0.575	0.681	0.535	0.589	0.332	0.549	0.005	0.549	0.384	0.552	0.570	0.491	0.181	0.186	0.599
18-year terms	0.190	0.118	0.001	0.570	0.575	0.681	0.535	0.589	0.321	0.506	0.007	0.549	0.384	0.561	0.417	0.497	0.147	0.112	0.633
Written emergency reasons	0.196	0.119	0.001	0.570	0.575	0.681	0.535	0.589	0.322	0.506	0.007	0.549	0.384	0.561	0.391	0.497	0.149	0.128	0.649
Automatic merits follow-up	0.194	0.119	0.001	0.570	0.575	0.681	0.535	0.589	0.317	0.507	0.007	0.549	0.384	0.567	0.369	0.497	0.148	0.089	0.588
No emergency merits gap	0.197	0.118	0.001	0.570	0.575	0.681	0.535	0.589	0.317	0.507	0.007	0.549	0.384	0.567	0.349	0.509	0.146	0.088	0.605
Strong recusal enforcement	0.194	0.118	0.001	0.570	0.575	0.681	0.535	0.589	0.321	0.506	0.007	0.549	0.384	0.561	0.418	0.513	0.148	0.113	0.632
Random merits panels	0.193	0.118	0.001	0.570	0.575	0.681	0.535	0.589	0.321	0.476	0.007	0.549	0.384	0.561	0.418	0.604	0.146	0.114	0.639
Emergency integrity package	0.196	0.118	0.001	0.570	0.575	0.681	0.535	0.589	0.317	0.507	0.006	0.549	0.384	0.567	0.369	0.513	0.146	0.089	0.598
Jurisdiction-stripping constraints	0.195	0.117	0.001	0.570	0.575	0.681	0.535	0.589	0.321	0.456	0.007	0.549	0.384	0.591	0.417	0.497	0.161	0.113	0.633
Legislative override window	0.194	0.119	0.001	0.570	0.575	0.681	0.535	0.589	0.322	0.507	0.007	0.549	0.384	0.563	0.417	0.497	0.156	0.113	0.633
Constitutional remand	0.201	0.118	0.001	0.570	0.575	0.681	0.535	0.589	0.286	0.465	0.007	0.549	0.384	0.580	0.418	0.501	0.145	0.113	0.698
Public-interest filter	0.200	0.117	0.001	0.570	0.575	0.681	0.535	0.589	0.321	0.421	0.007	0.549	0.384	0.562	0.417	0.501	0.147	0.112	0.643
Remand + override window	0.197	0.119	0.001	0.570	0.575	0.681	0.535	0.589	0.286	0.465	0.006	0.549	0.384	0.580	0.391	0.501	0.146	0.128	0.711
Panel + jurisdiction safeguards	0.197	0.118	0.001	0.570	0.575	0.681	0.535	0.589	0.322	0.425	0.006	0.549	0.384	0.589	0.402	0.631	0.156	0.131	0.669

Table 5 is the audit table for the concern that formal court design might overweight the model. It shows that the reported outputs now expose litigation-pipeline and downstream-enforcement quantities directly: certiorari admission, court-requested response, CVSG request, bar capital, claim strength, vehicle quality, genuine lower-court split rate, lower-court split depth, lower-court resistance, forum shopping, pre-review settlement, strategic plaintiff selection, repeat-player advantage, enforcement capacity, emergency opportunism, recusal incentive pressure, government noncompliance, emergency downstream effect, and precedent durability. The claimant, bar-capital, claim-strength, vehicle-quality, and genuine-split columns are shared-world composition diagnostics; the later columns show how institutional rules alter the realized pipeline and implementation environment. Those values are synthetic diagnostics, but they force the manuscript to show whether a formal design improvement is actually operating through intake, lower-court settlement, compliance, or emergency-order incentives.

Table 6 reports named-prior uncertainty bands from the parameter sweep. The overlap among the leading designs is the main point. The model is more informative about tradeoff patterns, such as emergency-restraint reducing shadow abuse and dual-court designs increasing conflict under the current disagreement rule, than about fine ordinal rankings among close scores.

Table 6: Campaign scores and uncertainty bands for selected designs

Scenario	Campaign score	Score 5/50/95	Shadow 5/50/95	Conflict 5/50/95
18-year terms	0.565	0.515/0.606/0.621	0.059/0.078/0.160	0.338/0.380/0.590
Written emergency reasons	0.562	0.509/0.602/0.618	0.032/0.049/0.123	0.336/0.381/0.591
Emergency merits follow-up	0.560	0.503/0.605/0.620	0.012/0.017/0.053	0.338/0.375/0.596
No emergency merits gap	0.566	0.512/0.609/0.625	0.007/0.011/0.043	0.339/0.373/0.593
Current-like	0.552	0.499/0.590/0.608	0.181/0.239/0.420	0.362/0.408/0.640
Strong recusal enforcement	0.559	0.509/0.599/0.616	0.057/0.077/0.163	0.341/0.381/0.590
Constitutional remand	0.566	0.514/0.601/0.617	0.056/0.077/0.164	0.337/0.380/0.591

Table 7: Sampled-prior uncertainty bands for selected designs

Scenario	Score	Rights	Shadow	Emerg. down.	Lower ct.	Gov. noncomp.	Conflict	Interpretation
No emergency relief without merits review	0.597/0.615/0.630	0.638/0.648/0.656	0.005/0.009/0.016	0.034/0.058/0.108	0.585/0.606/0.624	0.040/0.065/0.099	0.313/0.358/0.414	front-line cluster
18-year staggered terms + regular appointments	0.589/0.611/0.628	0.630/0.640/0.652	0.051/0.072/0.109	0.043/0.071/0.127	0.575/0.597/0.616	0.037/0.072/0.098	0.317/0.361/0.418	front-line cluster
Automatic merits follow-up for emergency relief	0.591/0.610/0.625	0.634/0.647/0.656	0.009/0.015/0.025	0.041/0.066/0.126	0.583/0.604/0.622	0.044/0.064/0.109	0.314/0.366/0.419	front-line cluster
Mandatory written emergency reasoning	0.586/0.608/0.623	0.626/0.639/0.649	0.029/0.045/0.073	0.049/0.075/0.139	0.574/0.596/0.616	0.047/0.070/0.102	0.320/0.369/0.421	front-line cluster
Constitutional remand before invalidation	0.584/0.607/0.623	0.626/0.640/0.647	0.050/0.074/0.112	0.043/0.071/0.129	0.587/0.609/0.628	0.042/0.066/0.107	0.317/0.368/0.416	front-line cluster
Randomized merits panels with en banc correction	0.584/0.606/0.622	0.631/0.641/0.650	0.052/0.073/0.109	0.044/0.070/0.126	0.575/0.598/0.617	0.042/0.066/0.105	0.319/0.365/0.412	front-line cluster
Independent recusal enforcement with substitutes	0.583/0.606/0.624	0.633/0.641/0.652	0.053/0.074/0.111	0.045/0.071/0.131	0.574/0.598/0.619	0.044/0.069/0.105	0.318/0.365/0.422	front-line cluster
Stylized current U.S.-like supreme court	0.568/0.594/0.616	0.627/0.645/0.658	0.167/0.234/0.324	0.079/0.122/0.196	0.543/0.576/0.602	0.051/0.095/0.150	0.337/0.401/0.462	overlapping uncertainty band

Table 7 adds a stricter uncertainty check: instead of only changing named assumption cases, it samples prior distributions over polarization, appointment capture, public pressure, emergency share, justice-pool size, and legislative-output profile components. The table makes close-score caution operational. If two designs’ sampled-prior bands overlap, the manuscript treats the comparison as a mechanism contrast rather than an ordinal ranking.

Table 8 states the most important interpretation check directly. If the top cluster changes only under high-emergency or high-conflict priors, the paper should describe that result as emergency- or conflict-contingent. If a design remains in the leading cluster but carries a rights, compliance, lower-court-resistance, or emergency-power caveat, the design recommendation should be weakened rather than reported as a clean improvement.

Figure 3 makes the emergency-procedure result visible. Similar rights and stability averages can coexist with very different emergency profiles, so the manuscript reports shadow-docket abuse, emergency legitimacy risk, and public confidence separately.

10 Mechanism-Level Summary

The scenario table names institutional bundles. Table 9 reports a second view: weighted paired contrasts from the mechanism-ablation report. Each row compares a current-like baseline to one institutional change across the same assumption cases. These are not causal estimates from observed courts, and they are not a substitute for the scenario table. They are a compact check on whether the named mechanisms move the outputs in the directions implied by the expectations.

Table 8: Sensitivity drivers and interpretation caveats

Prior	Top-cluster design	Score	Rights	Shadow	Emerg. down.	Gov. noncomp.	LC resist.	Caveat
Baseline institutional prior	No emergency relief without merits review	0.610	0.646	0.011	0.055	0.071	0.244	front-line cluster
Baseline institutional prior	60 percent invalidation threshold	0.608	0.633	0.064	0.079	0.070	0.248	rights-protection caveat
High emergency-share prior	18-year staggered terms + regular appointments	0.519	0.620	0.180	0.222	0.225	0.403	compliance caveat
High emergency-share prior	Public-interest litigation filter	0.519	0.630	0.175	0.215	0.218	0.402	compliance caveat
High rights-risk prior	18-year staggered terms + regular appointments	0.527	0.566	0.119	0.135	0.202	0.401	compliance caveat
High rights-risk prior	No emergency relief without merits review	0.527	0.600	0.026	0.098	0.226	0.398	compliance caveat
Weak democratic-mandate prior	60 percent invalidation threshold	0.546	0.604	0.094	0.125	0.180	0.377	compliance caveat
Weak democratic-mandate prior	Constitutional remand before invalidation	0.542	0.616	0.110	0.113	0.176	0.338	compliance caveat
High constitutional-conflict prior	60 percent invalidation threshold	0.512	0.601	0.142	0.194	0.235	0.439	compliance caveat
High constitutional-conflict prior	Constitutional remand before invalidation	0.509	0.614	0.153	0.166	0.235	0.400	compliance caveat
Imported legislative-family prior	No emergency relief without merits review	0.619	0.652	0.009	0.054	0.054	0.229	front-line cluster
Imported legislative-family prior	60 percent invalidation threshold	0.614	0.641	0.060	0.073	0.066	0.234	rights-protection caveat




Figure 3: Emergency-docket and public-confidence profile for selected institutional designs. The comparison highlights why emergency procedure is modeled directly: designs that are similar on rights protection and legal stability can differ sharply on shadow-docket abuse, emergency legitimacy risk, and public confidence.

The contrast table reinforces the main interpretation. Emergency restraint, written emergency reasoning, automatic merits follow-up, and strong recusal enforcement move shadow-docket and downstream-emergency diagnostics more directly than appointment or size reforms. Randomized panels, constitutional remand, public-interest filtering, jurisdiction-stripping constraints, and legislative override windows have more mixed profiles because they affect access, correction, compliance, or post-decision response rather than simply reducing court activity. These are synthetic mechanism

Table 9: Mechanism-level paired contrasts against the current-like design

Mechanism	Score	Rights	Shadow	Lower ct.	LC resist.	Enforce.	Gov. noncomp.	Emerg. opp.	Emerg. down.	Prec. dur.	Cost
Emergency restraint	+0.015	-0.002	-0.308	+0.031	-0.016	+0.017		-0.035	-0.225	-0.111	+0.011 +0.080
Mandatory written emergency reasoning	+0.010	-0.024	-0.249	+0.023	-0.011	+0.009		-0.039	-0.183	-0.064	+0.067 +0.080
Automatic merits follow-up	+0.010	-0.003	-0.300	+0.031	-0.017	+0.016		-0.039	-0.205	-0.111	-0.005 +0.096
Strong recusal enforcement	+0.007	-0.015	-0.210	+0.021	-0.011	+0.010		-0.035	-0.155	-0.078	+0.043 +0.087
Randomized merits panels	-0.005	+0.000	+0.000	+0.001	-0.000	+0.000		-0.003	+0.000	+0.001	+0.009 +0.045
Constitutional remand	+0.004	-0.002	-0.000	+0.020	-0.036	+0.019		-0.005	-0.000	-0.001	+0.064 +0.091
Public-interest litigation filter	+0.001	+0.005	-0.001	+0.001	-0.001	+0.000		-0.002	+0.000	-0.001	+0.017 +0.025
Legislative override window	-0.008	+0.001	-0.000	-0.004	+0.001	+0.002		+0.013	+0.000	+0.001	-0.004 +0.030
Jurisdiction-stripping constraints	-0.001	+0.003	-0.000	+0.001	-0.000	+0.030		+0.011	+0.000	-0.000	+0.002 +0.022
Emergency integrity bundle	+0.009	-0.003	-0.302	+0.033	-0.017	+0.017		-0.043	-0.206	-0.113	+0.009 +0.119
Remand and override-window bundle	-0.003	-0.005	-0.030	+0.016	-0.035	+0.020		+0.004	-0.026	+0.018	+0.079 +0.147
Panel and jurisdiction safeguards	-0.008	-0.008	-0.007	-0.002	+0.001	+0.028		+0.010	-0.014	+0.022	+0.035 +0.071
Council with concrete-review backstop	-0.008	+0.003	-0.002	+0.005	-0.001	+0.046		+0.012	+0.000	-0.002	+0.024 +0.098

contrasts, not observed causal estimates. Their use is to identify which mechanisms reduce arbitrary emergency power while preserving rights protection and which mechanisms merely move costs into noncompliance, administrative burden, or lower-court instability.

11 Limitations

The model is strongest as a disciplined comparison of institutional mechanisms, not as a predictive empirical model. It does not estimate causal effects from observed constitutional systems. It uses source-backed ranges, comparative institutional coding, and deterministic scenario campaigns to identify where rule bundles create tradeoffs that would otherwise be difficult to see. The strongest next validation steps are to automate raw-data refresh for SCDB, Harvard Law Review statistics, shadow-docket data, and recusal data; convert named priors into sampled prior distributions; add actor learning across repeated outcomes; calibrate review-path-specific admission denominators against raw certiorari, amparo, constitutional complaint, QPC, abstract-review, and concrete-referral data; and benchmark the legislative import contract against raw legislative simulator runs rather than only aggregate campaign CSVs.

A Institutional Parameter Provenance

Appendix Table 11 records how representative metrics should be read. The point is not to create another score table; it is to make the denominator or scale explicit before any empirical, synthetic, or normative inference is drawn.

B Model Mechanics

This appendix records the main state machine and scoring formulas so the manuscript does not treat the simulator as a black box. Each run draws a filed universe from a shared world seed. For each scenario, the same filed matters are reviewed by a court process with a scenario-specific design and a scenario-specific random stream. Let i index filed matters and s index institutional scenarios.

Admission begins with a path-specific petition probability. In the compact notation below, c is legal conflict, a is public attention, g is SG signal, d is split depth, k is bar capital, m is claim strength, r is repeat-player advantage, p is strategic-plaintiff selection, f is forum-shopping pressure, q is settlement pressure, and v is vehicle-defect risk:

$$P_{is} = \text{clamp}_{[0,1]}(b_{\text{petition}} + .08c_i + .08a_i + .06g_i + .08d_i \\ + .05k_i + .05m_i + .06r_i + .08p_i + .04f_i - .05q_i - .22v_i).$$

Conditional on a petition or referral, the admission score combines certiorari pressure, lower-court conflict, lower-court error, lower-court split depth, genuine split status, SG signal, court-requested response, CVSG request, amicus intensity, split maturity, relist signal, claimant type, bar capital, claim strength, vehicle quality, specialist counsel, repeat-player advantage, strategic plaintiff selection, forum-shopping pressure, settlement pressure, conditional reversal probability, access-path adjustment, institutional gatekeeping, vehicle defects, and vacancy deadlock. Merits transfer is then path-specific: emergency applications transfer through the application’s merits-follow-through probability, complaints and amparo transfer at a lower rate, and ordinary referrals or cert-like paths transfer at a higher baseline reduced by settlement pressure.

After admission, the court process chooses strategic-actor response, emergency status, pre-review settlement, initial merits review, panel or en banc routing, recusals, quorum, votes, emergency order, invalidation, auxiliary-review outcome, override attempt, and post-decision implementation response. The direct metrics are case averages over the filed universe. For example, rights-claimant success is one when a high-rights-burden claimant receives invalidation or shadow relief and partial when settlement or constitutional remand preserves a claimant’s route without immediate invalidation;

Table 10: Illustrative provenance map for modeled institutional features

Feature family	Model representation	Provenance status
Appointment and tenure	Appointment method, confirmation threshold, vacancy deadlock risk, term renewability, retirement age, replacement timing	Comparative institutional coding and scenario assumptions
Review access	Access path, petition type, court-requested response, CVSG request, admission score, claimant type, bar capital, claim strength, vehicle quality, lower-court split depth, strategic plaintiff selection, repeat-player advantage, merits transfer, screen-out	Normalized source-register rows and modeled admission rules
Emergency procedure	Application class, relief request, response request, full-court referral, reason giving, status-quo effect, merits follow-through, downstream effect	Shadow-docket source tables plus scenario rules
Decision rule	Voting threshold, panel/en banc routing, quorum failure, recusal consequence, per-remedy thresholds	Court-design coding and institutional assumptions
Post-decision politics	Override attempt, rights carveout, formal response, practical implementation, government noncompliance, open defiance	Comparative public-law coding and strategic actor assumptions
Outputs	Stability, rights protection, claimant success, legitimacy, conflict, responsiveness, precedent durability, lower-court compliance	Simulated outcomes with calibration guardrails

public confidence is legitimacy minus penalties for emergency legitimacy risk, partisan alignment, appointment-pressure campaigns, and court-curbing response; constitutional conflict rises with conflict potential, executive defiance, lower-court conflict, lower-court resistance, accumulated state conflict, government noncompliance, emergency downstream effect, invalidating high-mandate laws, shadow relief, cross-court disagreement, override attempts, and repeated override loops.

The directional score is intentionally secondary:

$$\text{DirectionalScore}_s = \text{mean}(\text{StabilityRights}_s, \text{LegitimacyControl}_s, \text{ClaimantSuccess}_s, \text{PrecedentDurability}_s, \text{LowerCourtCompliance}_s, \text{EliteAcceptance}_s, 1 - \text{AdministrativeCost}_s).$$

The stability-rights component averages legal stability, precedent durability, lower-court compliance, enforcement capacity, lower-court-resistance restraint, rights protection, reversal restraint, and

conflict restraint. The legitimacy-control component averages legitimacy, democratic responsiveness, independence-accountability balance, partisan-alignment restraint, shadow-abuse restraint, emergency-legitimacy-risk restraint, emergency-downstream restraint, government-noncompliance restraint, recusal-pressure restraint, forum-shopping restraint, emergency-opportunism restraint, and strategic-pressure restraint. Because those choices are normative, the paper reports the direct outputs separately and uses the directional score only as a reading aid. Score differences below .010 are treated as indistinguishable in the manuscript’s substantive interpretation.

Data Availability Statement

The simulator, normalized calibration tables, frozen legislative-output fixtures, deterministic seeds, scenario definitions, generated reports, figure-generation scripts, standalone figure exports, source-audit file (`paper/source-audit.csv`), source-provenance manifests, and report manifests are maintained as a reproducible project. Public repository identifiers are withheld in this anonymous review manuscript. The intended replication package includes source code, build instructions, normalized input data, generated campaign artifacts, generated figure fragments, standalone figure files, manifest hashes, and documentation distinguishing observed inputs, coding choices, model assumptions, and simulated outputs. For final publication, the package should be deposited in the Journal of Law and Courts Dataverse or another repository linked from the JLC Dataverse rather than being available only on request.

Funding Statement

No external funding is reported for this version of the manuscript.

Competing Interests

Competing interests: The author(s) declare none.

AI Assistance Statement

The author(s) used OpenAI ChatGPT Deep Research and OpenAI Codex during May 2026, accessed through OpenAI’s ChatGPT and Codex interfaces, to assist with research-prompt drafting, source-synthesis review, code generation, manuscript drafting, formatting checks, and local build verification. The author(s) did not list AI systems as authors, did not train or modify these systems, and reviewed, edited, and are responsible for the manuscript, simulator code, calibration choices, generated reports, generated figures, references, and all claims. Outputs from AI-assisted research were treated as leads or author-created synthesis notes; third-party sources, datasets, and claims were checked against cited sources or local project artifacts before inclusion. No AI-generated visual content appears in the manuscript; all figures are generated deterministically from project CSV outputs by repository code.

References

Baude, W. (2015). Foreword: The supreme court’s shadow docket. *New York University Journal of Law and Liberty*, 9:1–47.

- Black, R. C. and Epstein, L. (2005). Recusals and the “problem” of an equally divided supreme court. *Journal of Appellate Practice and Process*, 7(1):75–99. Article and data page available at <https://epstein.wustl.edu/recusal>.
- Brewer-Carías, A. R. (2009). *Constitutional Protection of Human Rights in Latin America: A Comparative Study of Amparo Proceedings*. Cambridge University Press, Cambridge.
- Clayton, C. W. and Gillman, H., editors (1999). *Supreme Court Decision-Making: New Institutionalalist Approaches*. University of Chicago Press, Chicago.
- Comella, V. F. (2009). *Constitutional Courts and Democratic Values: A European Perspective*. Yale University Press, New Haven.
- Dixon, R. and Ginsburg, T. (2017). The forms and limits of constitutions as political insurance. *International Journal of Constitutional Law*, 15(4):988–1012.
- Epstein, L. and Knight, J. (1998). *The Choices Justices Make*. CQ Press, Washington, DC.
- Gardbaum, S. (2013). *The New Commonwealth Model of Constitutionalism: Theory and Practice*. Cambridge University Press, Cambridge.
- Garoupa, N. and Ginsburg, T. (2009). Guarding the guardians: Judicial councils and judicial independence. *American Journal of Comparative Law*, 57(1):103–134.
- Ginsburg, T. (2003). *Judicial Review in New Democracies: Constitutional Courts in Asian Cases*. Cambridge University Press, Cambridge.
- Harvard Law Review (2025). The supreme court, 2024 term – the statistics. 139 Harvard Law Review 430. Subject-matter, merits, emergency-relief, and voting-alignment tables. Available at <https://harvardlawreview.org/category/statistics/>. Accessed May 5, 2026.
- Hirschl, R. (2004). *Towards Juristocracy: The Origins and Consequences of the New Constitutionalism*. Harvard University Press, Cambridge, MA.
- Hogg, P. W. and Bushell, A. A. (1997). The charter dialogue between courts and legislatures. *Osgoode Hall Law Journal*, 35(1):75–124.
- Kastellec, J. P. and Taboni, A. R. (2026). A database of the united states supreme court’s shadow docket, 1993–2025. *Journal of Law and Courts*, pages 1–18. Database available at <https://www.shadowdocketdata.com/data>.
- Kommers, D. P. and Miller, R. A. (2012). *The Constitutional Jurisprudence of the Federal Republic of Germany*. Duke University Press, Durham, NC, 3 edition.
- Segal, J. A. and Spaeth, H. J. (2002). *The Supreme Court and the Attitudinal Model Revisited*. Cambridge University Press, Cambridge.
- Staton, J. K. (2010). *Judicial Power and Strategic Communication in Mexico*. Cambridge University Press, Cambridge.
- Supreme Court Database (2025). Online codebook and modern database 2025 release 01. Issue-area, declaration-of-unconstitutionality, and formal-alteration-of-precedent variables. Available at <https://scdb.la.psu.edu/online-codebook/> and <https://scdb.la.psu.edu/data/2025-release-01/>. Accessed May 5, 2026.

- Sweet, A. S. (2000). *Governing with Judges: Constitutional Politics in Europe*. Oxford University Press, Oxford.
- Tushnet, M. (2008). *Weak Courts, Strong Rights: Judicial Review and Social Welfare Rights in Comparative Constitutional Law*. Princeton University Press, Princeton.
- Vanberg, G. (2005). *The Politics of Constitutional Review in Germany*. Cambridge University Press, Cambridge.
- Vladeck, S. I. (2023). *The Shadow Docket: How the Supreme Court Uses Stealth Rulings to Amass Power and Undermine the Republic*. Basic Books, New York.

Table 11: Metric semantics and manuscript use

Family	Metric	Use	Compatibility	Manuscript interpretation
certiorari	paid cert petition share	loose calibration	near pathway match	conditional paid/IFP split is closer than all-case paidPetitionRate, but the simulator still excludes many non-review docketed matters
certiorari	certiorari admission rate	proxy context	denominator mismatch	the source includes all docketed matters; the simulator denominator is already a constitutional-review cert-path subset
emergency	emergency grant conditional rate	loose calibration	conditional near match	conditional rate is closer than all-case emergencyGrantRate, but the simulated emergency universe is not limited to noncapital applications
complaint referral	admission rate	proxy context	denominator mismatch	this is a strong complaint-filter source but not a whole-universe admission target
lower court compliance	lower court compliance	proxy context	mechanism proxy	the source anchors direction and magnitude of alignment pressure, not a compliance-level target
lower court compliance	government noncompliance rate	loose calibration	mechanism proxy	closest available implementation anchor, but still a study-sample proxy rather than a court-wide noncompliance rate
override remand	override rate	design prior	design prior	duration informs institutional design, not observed behavioral override success
headline	directional score	reading aid	not empirical target	used only for ordering and clustering; close score differences are not substantive rankings
headline	rights claimant success	synthetic output	not empirical target	direct model output, not blended into empirical validation
headline	public confidence	synthetic output	not empirical target	comparative index for mechanism direction, not a survey-estimated confidence level
headline	constitutional conflict	synthetic output	not empirical target	mechanism summary metric; individual channels remain reported separately

Table 12: Representative model weights and decision rules

Component	Representative rule	Interpretation
Admission	0.18 + .22 cert pressure +.12 lower-court conflict +.10 lower-court error +.13 split depth +.06 genuine split +.15 SG signal +.10 amicus signal +.08 split maturity +.07 relist signal +.06 bar capital +.10 claim strength +.09 vehicle quality +.05 specialist counsel +.05 repeat-player signal +.05 strategic-plaintiff signal +.02 forum-shopping signal +.08 conditional reversal -.06 settlement pressure -.16 vehicle defect, plus path adjustment, CFR, and CVSG increments	Admission is an institutional screen, not a fixed docket quota.
Voting	Legal concern receives weight .42; rights concern uses rights burden, justice rights sensitivity, and weight .44; claim strength, vehicle quality, and bar capital can modestly increase merits credibility; democratic mandate, legislative quality, institutionalism, emergency deference, repeat-player pressure, forum-shopping signal, precedent stability, and accountability pull reduce or increase invalidation propensity.	Merits votes mix legal, rights, strategic, institutional, and emergency-deference considerations.
Emergency abuse	Open emergency rule base .32, no-merits penalty .22, shadow stay .24, written reasons -.10, automatic merits follow-up -.12, merits acceleration -.07, partisan salience .12, ambiguity .08, emergency opportunism .10.	Shadow-docket abuse is driven by opacity, opportunistic emergency filing, and merits-like relief without ordinary merits process.
Legitimacy	Base .48, public attention .16, legislative quality .10, merits reason-giving .18, recusal discipline up to .16, council warning .06, shadow abuse -.24, partisan alignment -.20.	Public legitimacy rewards visible process and penalizes opaque or partisan-looking intervention.
Compliance	Lower-court compliance rises with enforcement capacity and merits clarity, and falls with lower-court conflict, error risk, split depth, ideological drift, lower-court resistance, state conflict, invalidation shock, and government noncompliance.	Enforcement and implementation are modeled separately from formal legal validity.
Conflict	Conflict potential .38, executive defiance .18, lower-court conflict .08, lower-court resistance .10, forum shopping .04, conflict load .12, legislative defiance .10, government noncompliance, emergency downstream effect, invalidating high-mandate laws .14, shadow relief .12, cross-court disagreement .16, override attempt .06.	Constitutional conflict is cumulative and rises when legal disagreement becomes interbranch disagreement.
Override	Attempt pressure combines override pressure .34, democratic mandate .30, public attention .18, conflict load .12, legislative defiance .12, override adaptation .10, and rights-burden penalty -.24.	Accountability devices affect implementation only after invalidation and can generate rights carveouts or repeated override loops.